

SkillForge

Measure the gap. **Then close it.**

System design

you 2.0 · role 6.0



skillforge-eight-pink.vercel.app · live

Placement season is a triage problem, and nobody hands you the ruler.

- 01 A final-year student has a resume that **overstates some things and understates others** — and no way to know which.

- 02 Twenty things they could study; a fixed number of weeks. **Which one actually changes whether they get an offer?**

- 03 Career platforms answer with **generic recommendations**. AI chatbots answer with a different number every time you ask.

The model extracts. Code decides.

One model call reads the resume — not to score it, but to report **evidence**: what was built, what shipped, what has tests, with a page and line for every claim. Everything after that is arithmetic over a **published, versioned benchmark**.

1

model call per analysis — everything else is pure TypeScript

0–10

HOW IT WORKS

levels snapped to human-written rubric ladders, never generated

One pipeline, one checkpoint.

100%


Parse



navigate

PDF → page/line-indexed text. OCR fallback for scans.





Extract

The one model call: evidence signals + cited flags. No score field exists.




Verify

Every quote checked against its cited page & line. Hallucinations dropped.



Score · Rank · Schedule

Weighted gap points; a 14-week plan from a topological sort.



THE PRODUCT



From one upload: the map, the plan,

and something to do this week.

The skill map

Solid is what the resume proves. The notch is what the role demands. The hatch is the work.


System design



Docker & CI/CD



DSA – graphs & DP



12 tracks measured · 5 gaps open

readiness **71** / 100

The plan

14 weeks, three lanes, prerequisite-ordered — Docker lands in week 1 *because* the load test needs it — packed under the hours the student actually has.

The practice

Projects and interview drills ranked by **weighted gap points**; certifications re-scored live in the browser as a budget slider moves.

The proof

"Why not just paste it into ChatGPT?"

"You can ask an AI what it thinks of your resume. **You can't ask it to measure.**"

LIVE DEMO · 3 MINUTES

Watch the numbers earn themselves.

```
:15 The console boots — ask the mentor on the landing page what  
this is
```

:45 Upload a resume, narrate the pipeline while the checklist runs

:30 Skill map — one gauge, its rubric rung, **its resume citation**

:20 **Switch target role** — everything re-measures instantly

```
:20 /benchmarks — the ruler, logged out
```

:50 Mentor: "why is my system design a 2?" · Studio: "show my Docker gap" → **the answer renders as a gauge**

Shipped — not scaffolded.

\$0

total model spend — keyless free tier, and an architecture with almost nothing to pay for

81

unit tests on the deterministic core, plus 11 `check:*` scripts against live services

6

agents, each with a hard structural boundary on what it cannot do

Next.js 16 · Mastra · OpenCode Zen (keyless) · Neon Postgres + Drizzle ·
Clerk · Tailwind v4 · OpenUI Lang

skillforge-eight-pink.vercel.app — deployed, public benchmarks,
live upload.

Readiness 71 / 100 – and it knows why

